

Name: Muhammad AhadImran

ID: F23607034

PFAI ESE

QUESTION # 02

Part1 & 2: R Language

1)

Using **mtcars** dataset in R:

- Group data by number of gears
- For each group, calculate the average mpg and hp
- Create a grouped bar plot comparing both averages using Include code and output.

```
# Load required library
library(ggplot2)
library(dplyr)

# Group by gear and calculate average mpg and hp
mtcars_grouped <- mtcars %>%
  group_by(gear) %>%
  summarise(avg_mpg = mean(mpg), avg_hp = mean(hp))

# Reshape data for plotting
mtcars_long <- tidyr::pivot_longer(mtcars_grouped,
                                   cols = c(avg_mpg, avg_hp),
                                   names_to = "metric",
                                   values_to = "value")

# Create grouped bar plot
ggplot(mtcars_long, aes(x = factor(gear), y = value, fill = metric)) +
  geom_bar(stat = "identity", position = position_dodge(width = 0.8)) +
```

```
labs(title = "Average MPG and HP by Number of Gears",  
      x = "Number of Gears",  
      y = "Average Value",  
      fill = "Metric") +  
theme_minimal()
```



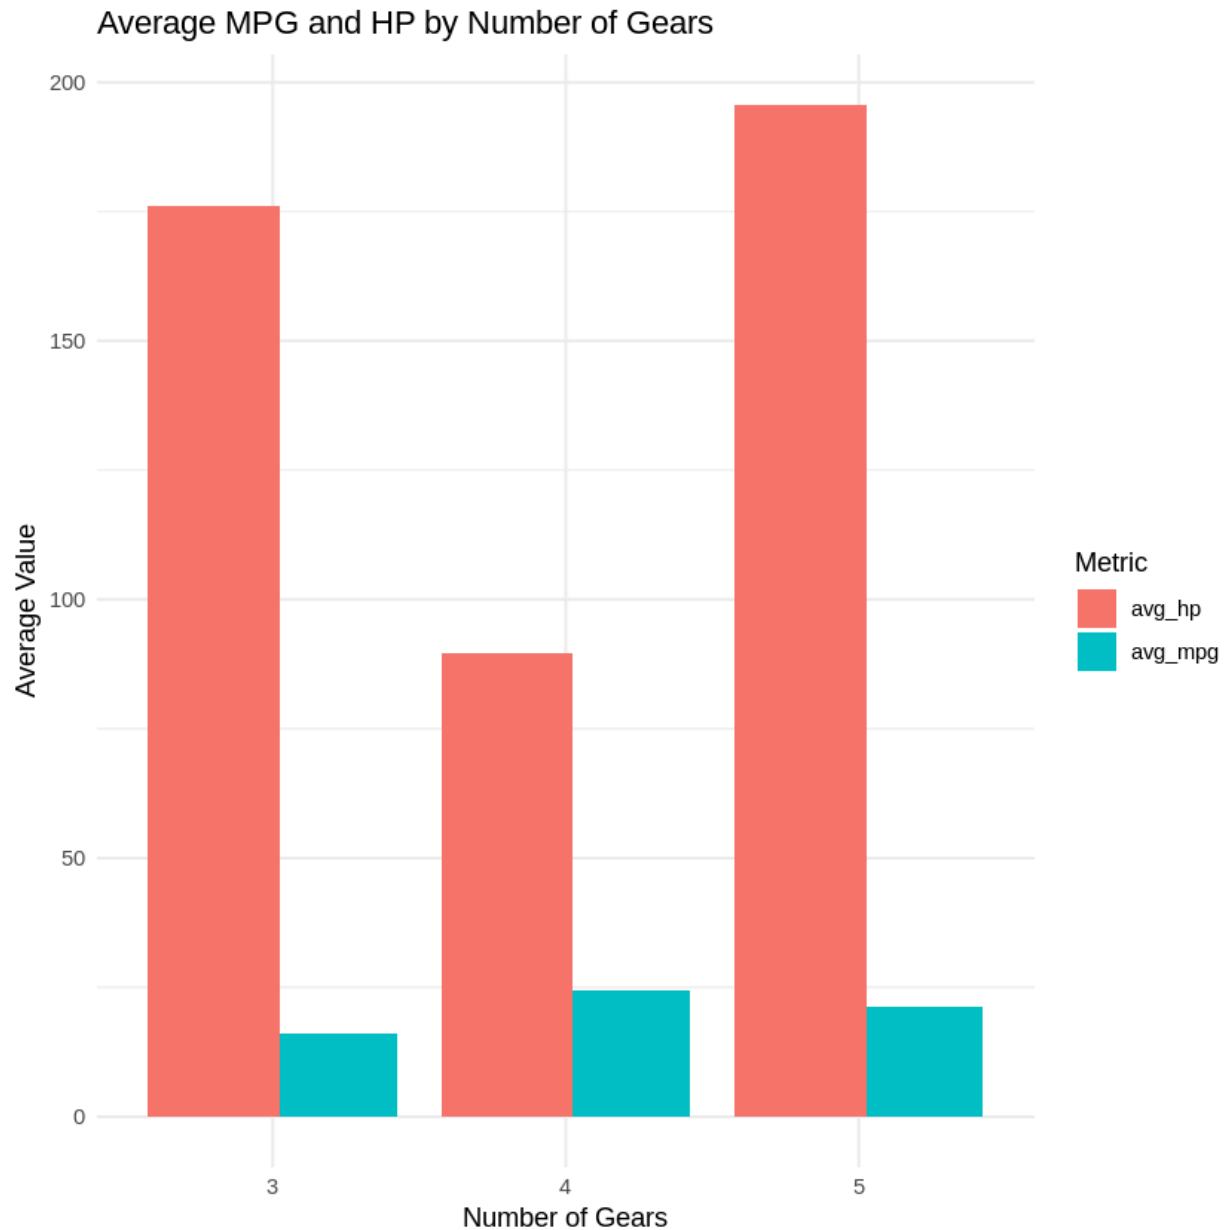
Attaching package: 'dplyr'

The following objects are masked from 'package:stats':

filter, lag

The following objects are masked from 'package:base':

intersect, setdiff, setequal, union



2)

Write a R script that:

- Splits `mtcars` dataset in train/test sets
- Scales the data
- Trains a Random Forest to predict mpg
- Evaluates performance using RMSE and R-squared
- Add comments explaining each step

```
# Load necessary libraries
install.packages("caret")
install.packages("randomForest")
library(caret)
library(randomForest)

# Set random seed for reproducibility
set.seed(123)

# Split the data into training and testing sets (80% train, 20%
test)
train_index <- createDataPartition(mtcars$mpg, p = 0.8, list =
FALSE)
train_data <- mtcars[train_index, ]
test_data <- mtcars[-train_index, ]

# Separate features (X) and target variable (y) for training and
testing sets
train_x <- train_data[, -which(names(train_data) == "mpg")]
train_y <- train_data$mpg
test_x <- test_data[, -which(names(test_data) == "mpg")]
test_y <- test_data$mpg

# Scale the numerical features using preProcess
# This is important for many machine learning algorithms, though
less critical for Random Forest
# We will fit the scaler on the training data and apply it to both
training and testing data
preproc_param <- preProcess(train_x, method = c("center", "scale"))
train_x_scaled <- predict(preproc_param, train_x)
test_x_scaled <- predict(preproc_param, test_x)

# Combine scaled features and target variable back into data frames
for training and testing
train_scaled <- cbind(train_x_scaled, mpg = train_y)
test_scaled <- cbind(test_x_scaled, mpg = test_y)

# Train a Random Forest model to predict mpg
```

```
# The formula specifies mpg as the target and all other columns as
predictors
# ntree specifies the number of trees in the forest
# mtry specifies the number of variables randomly sampled at each
split
rf_model <- randomForest(mpg ~ ., data = train_scaled, ntree = 500,
mtry = floor(ncol(train_scaled)/3))

# Make predictions on the scaled test data
predictions <- predict(rf_model, test_scaled)

# Evaluate the model performance

# Calculate Root Mean Squared Error (RMSE)
# RMSE measures the standard deviation of the residuals (prediction
errors)
rmse <- sqrt(mean((test_scaled$mpg - predictions)^2))
cat("RMSE:", rmse, "\n")

# Calculate R-squared
# R-squared represents the proportion of the variance for a
dependent variable that's explained by the independent variables in
a regression model.
# It ranges from 0 to 1, where 1 indicates a perfect fit.
sse <- sum((test_scaled$mpg - predictions)^2) # Sum of squared
errors
sst <- sum((test_scaled$mpg - mean(test_scaled$mpg))^2) # Total sum
of squares
r_squared <- 1 - (sse / sst)
cat("R-squared:", r_squared, "\n")
```

```
➔ Installing package into '/usr/local/lib/R/site-library'
(as 'lib' is unspecified)

Installing package into '/usr/local/lib/R/site-library'
(as 'lib' is unspecified)

randomForest 4.7-1.2

Type rfNews() to see new features/changes/bug fixes.

Attaching package: 'randomForest'

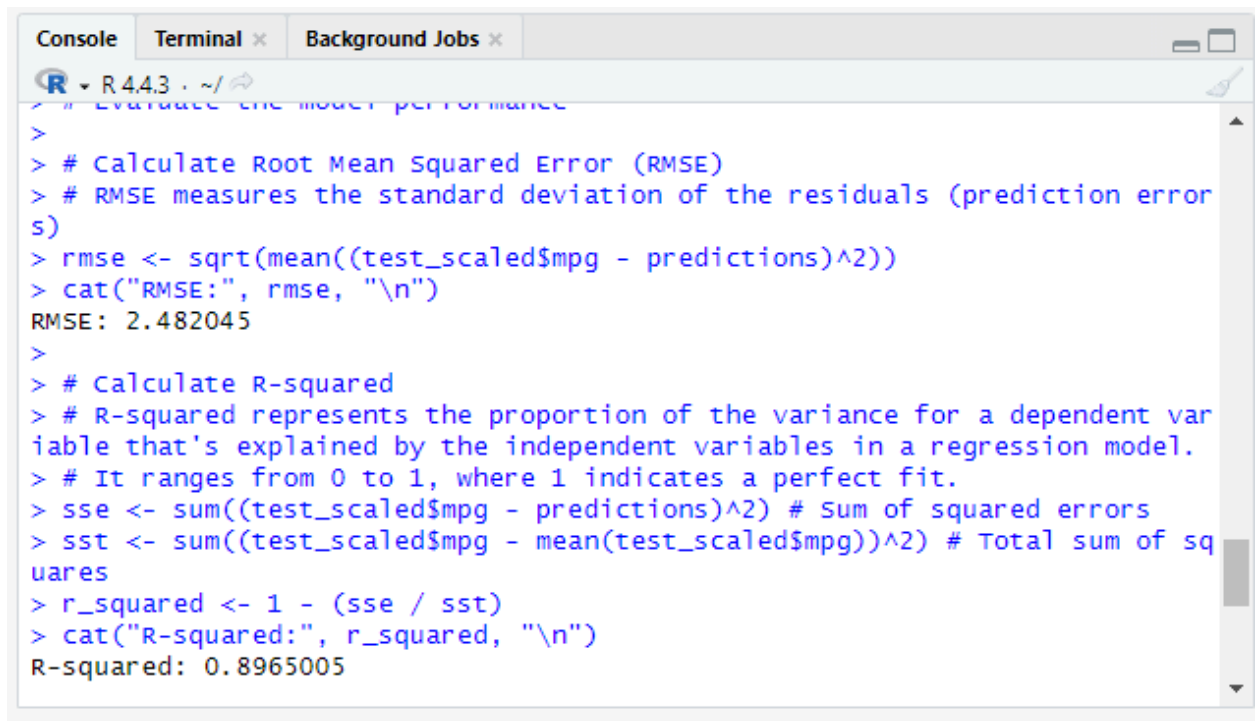
The following object is masked from 'package:dplyr':

  combine

The following object is masked from 'package:ggplot2':

  margin

RMSE: 2.482045
R-squared: 0.8965005
```



The screenshot shows an R console window with the following content:


```
Console Terminal x Background Jobs x
R 4.4.3 ~ /
> # Evaluate the model performance
>
> # Calculate Root Mean Squared Error (RMSE)
> # RMSE measures the standard deviation of the residuals (prediction error
s)
> rmse <- sqrt(mean((test_scaled$mpg - predictions)^2))
> cat("RMSE:", rmse, "\n")
RMSE: 2.482045
>
> # calculate R-squared
> # R-squared represents the proportion of the variance for a dependent var
iable that's explained by the independent variables in a regression model.
> # It ranges from 0 to 1, where 1 indicates a perfect fit.
> sse <- sum((test_scaled$mpg - predictions)^2) # sum of squared errors
> sst <- sum((test_scaled$mpg - mean(test_scaled$mpg))^2) # Total sum of sq
uares
> r_squared <- 1 - (sse / sst)
> cat("R-squared:", r_squared, "\n")
R-squared: 0.8965005
```

Part 3: Python Language

Given a dataset with missing values and inconsistent formatting, write a python script using the pandas library to clean the data. Your cscript should handle missing values using proper imputation techniques and standardize categorical variables for consistency.

Using the matplotlib and seaborn libraries, generate meaningful visualizations for a dataset containing multiple numeric variables. Create a heatmap to illustrate correlations, and use histograms or box plots to show the distribution of key feature.

```
import os
print(os.listdir('sample_data'))
```

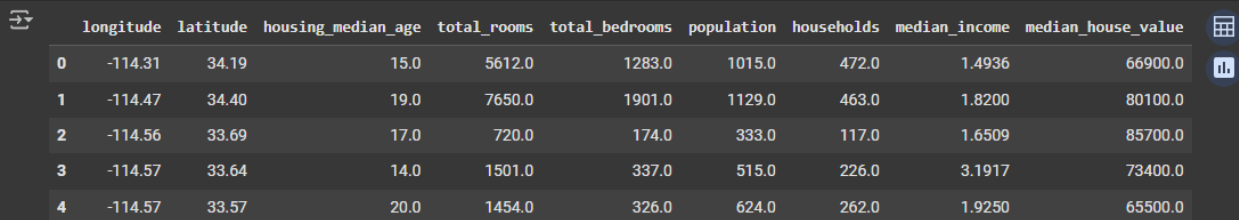


```
['anscombe.json', 'README.md', 'california_housing_train.csv', 'mnist_test.csv', 'california_housing_test.csv', 'mnist_train_small.csv']
```

```
import pandas as pd

df = pd.read_csv('sample_data/california_housing_train.csv')
```

```
display(df.head())
```



The image shows a Jupyter Notebook interface with a table of housing data. The table has 10 columns: longitude, latitude, housing_median_age, total_rooms, total_bedrooms, population, households, median_income, and median_house_value. The first five rows are displayed, indexed from 0 to 4. The interface includes a code editor on the left and a table viewer on the right.

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	median_house_value
0	-114.31	34.19	15.0	5612.0	1283.0	1015.0	472.0	1.4936	66900.0
1	-114.47	34.40	19.0	7650.0	1901.0	1129.0	463.0	1.8200	80100.0
2	-114.56	33.69	17.0	720.0	174.0	333.0	117.0	1.6509	85700.0
3	-114.57	33.64	14.0	1501.0	337.0	515.0	226.0	3.1917	73400.0
4	-114.57	33.57	20.0	1454.0	326.0	624.0	262.0	1.9250	65500.0

Handle missing values

```
display(df.isnull().sum())
```



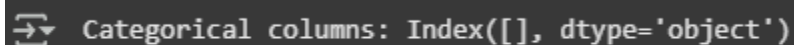
The image shows a Jupyter Notebook interface displaying the result of the `df.isnull().sum()` command. It shows a list of columns and their corresponding sum of missing values, all of which are 0. The dtype is int64.

	0
longitude	0
latitude	0
housing_median_age	0
total_rooms	0
total_bedrooms	0
population	0
households	0
median_income	0
median_house_value	0

dtype: int64

Standardize categorical variables

```
categorical_cols = df.select_dtypes(include=['object',  
'category']).columns  
print("Categorical columns:", categorical_cols)  
  
for col in categorical_cols:  
    print(f"\nUnique values in '{col}':")  
    display(df[col].value_counts())
```



The image shows a Jupyter Notebook interface displaying the output of the code. It shows the categorical columns as an Index with dtype='object'.

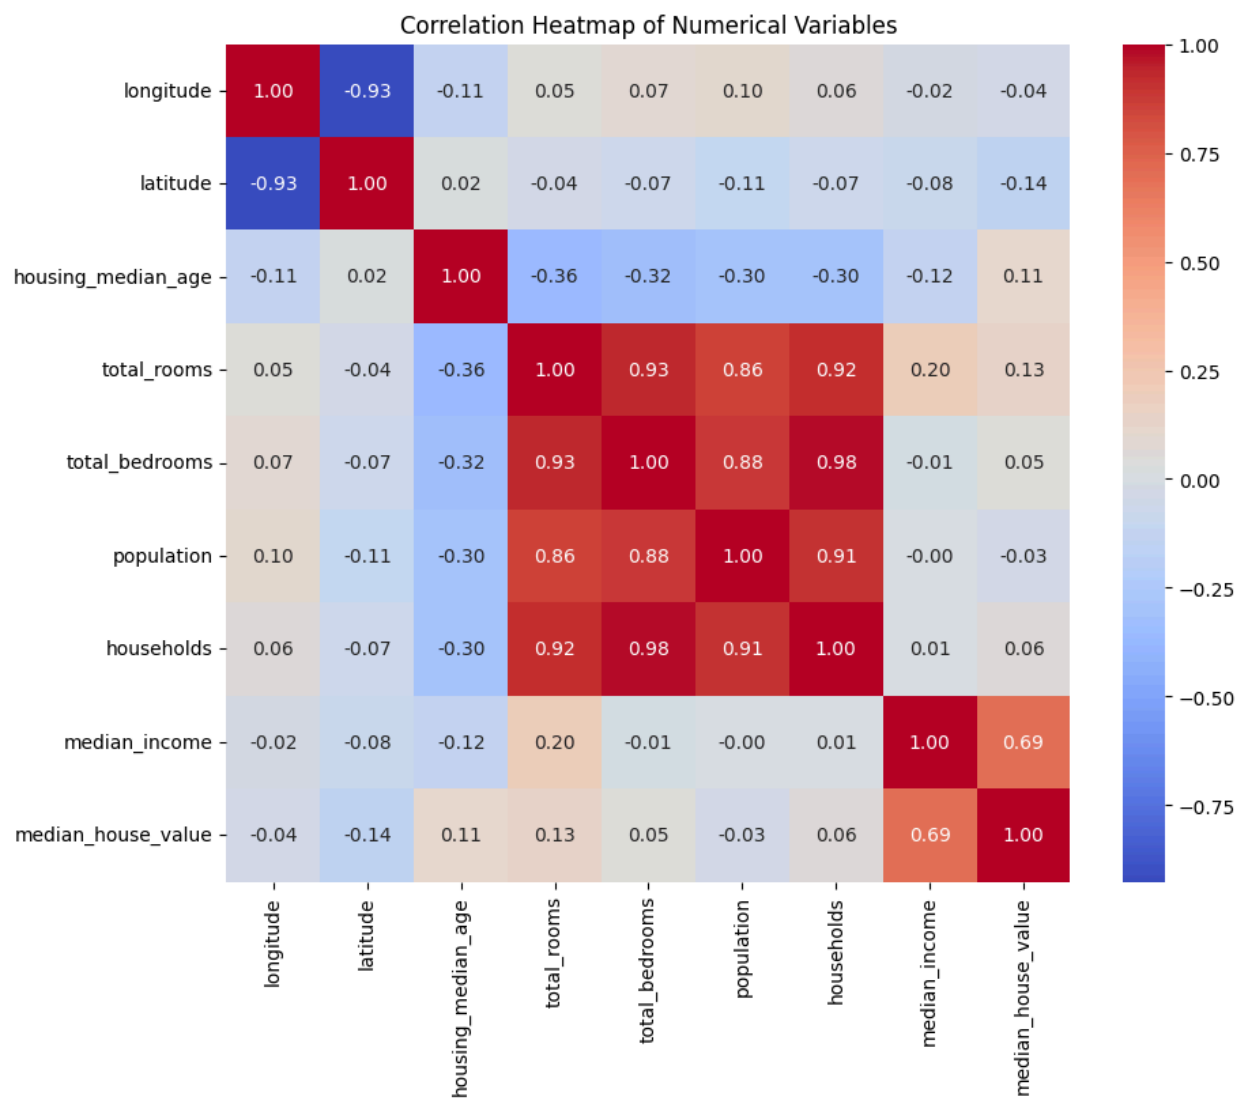
```
Categorical columns: Index([], dtype='object')
```


Visualize correlations

```
correlation_matrix = df.corr()

import seaborn as sns
import matplotlib.pyplot as plt

plt.figure(figsize=(10, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt=".2f")
plt.title('Correlation Heatmap of Numerical Variables')
plt.show()
```



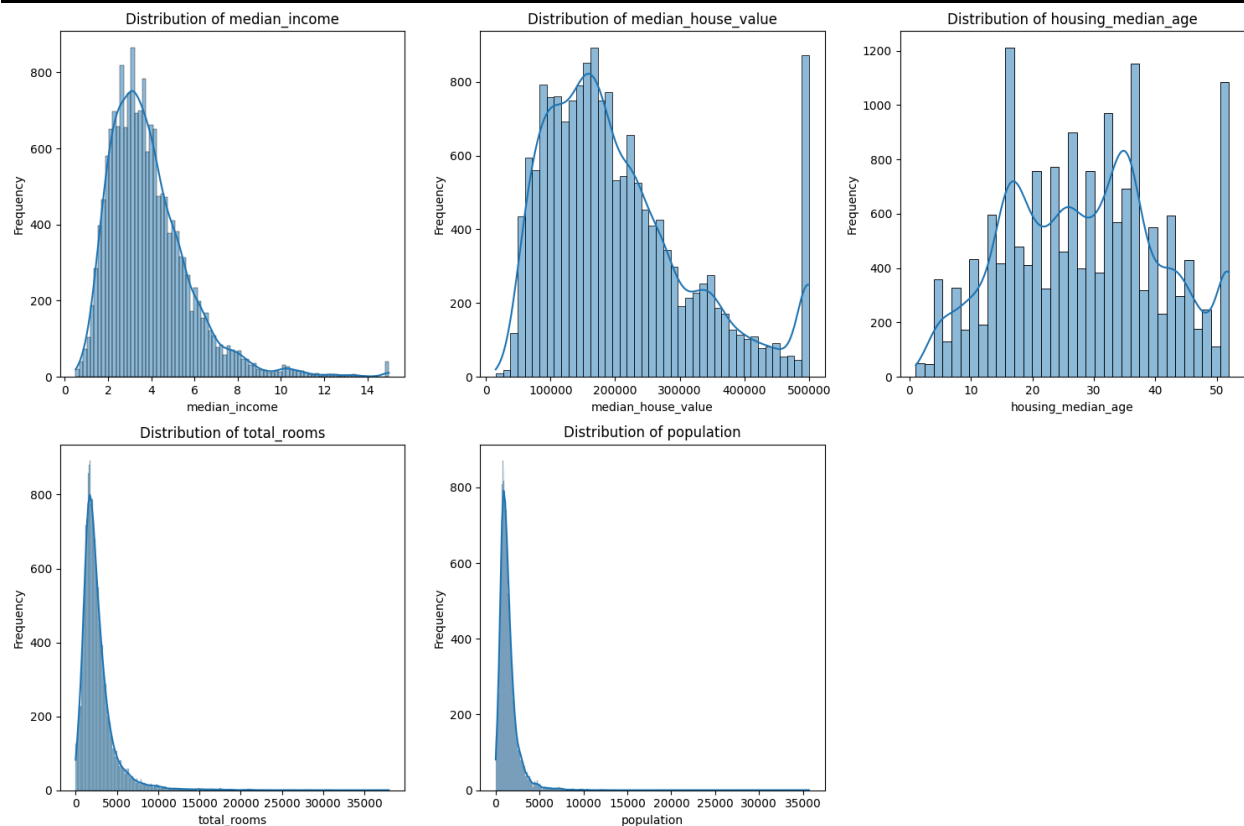
Visualize distributions

```
import matplotlib.pyplot as plt
import seaborn as sns

key_numerical_cols = ['median_income', 'median_house_value',
                      'housing_median_age', 'total_rooms', 'population']

plt.figure(figsize=(15, 10))
for i, col in enumerate(key_numerical_cols):
    plt.subplot(2, 3, i + 1)
    sns.histplot(df[col], kde=True)
    plt.title(f'Distribution of {col}')
    plt.xlabel(col)
    plt.ylabel('Frequency')

plt.tight_layout()
plt.show()
```



Question # 01

Part 1: GO LANGUAGE:

1)

```
package main

import (
    "fmt"
)

// BankAccount struct
type BankAccount struct {
    Owner    string
    Balance  float64
}

// Deposit method: adds amount to balance
func (b *BankAccount) Deposit(amount float64) {
    b.Balance += amount
    fmt.Printf("%s deposited: $%.2f\n", b.Owner, amount)
}

// Withdraw method: subtracts if enough balance
func (b *BankAccount) Withdraw(amount float64) {
    if amount > b.Balance {
        fmt.Printf(" Withdrawal of $%.2f failed: Insufficient
funds!\n", amount)
    } else {
        b.Balance -= amount
        fmt.Printf(" %s withdrew: $%.2f\n", b.Owner, amount)
    }
}

// BalanceInquiry method: shows current balance
func (b BankAccount) BalanceInquiry() {
```

```

        fmt.Printf(" Current balance for %s: $%.2f\n", b.Owner,
b.Balance)
    }

func main() {
    // Create a bank account
    account := BankAccount{Owner: "Ahad"}

    // 3 demo transactions
    account.Deposit(1000)
    account.Withdraw(250)
    account.Withdraw(1000) // This should fail

    // Final balance check
    account.BalanceInquiry()
}

```

```

Ahad deposited: $1000.00
Ahad withdrew: $250.00
Withdrawal of $1000.00 failed: Insufficient funds!
Current balance for Ahad: $750.00

Program exited.

```

2)

```

package main

import (
    "net/http"

    "github.com/gin-gonic/gin"
)

// Define a Student struct
type Student struct {
    Name string `json:"name"`
    Age  int    `json:"age"`
    GPA  float64 `json:"gpa"`
}

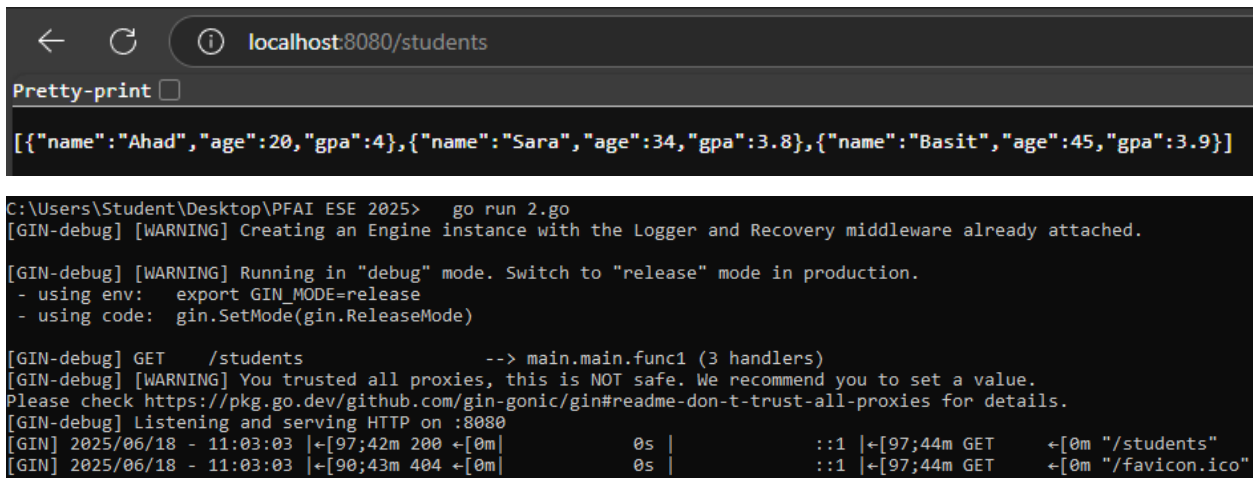
```

```
// Main function
func main() {
    // Initialize Gin router
    router := gin.Default()

    // Create GET /students endpoint
    router.GET("/students", func(c *gin.Context) {
        students := []Student{
            {Name: "Ahad", Age: 20, GPA: 4.0},
            {Name: "Sara", Age: 34, GPA: 3.8},
            {Name: "Basit", Age: 45, GPA: 3.9},
        }

        // Return JSON with status 200
        c.JSON(http.StatusOK, students)
    })

    // Start the server on port 8080
    router.Run(":8080")
}
```



The screenshot shows a web browser at `localhost:8080/students` displaying a JSON array of student data. Below the browser, a terminal window shows the command `go run 2.go` and the output of the application, including GIN debug messages and HTTP logs.

```

C:\Users\Student\Desktop\PFAI ESE 2025> go run 2.go
[GIN-debug] [WARNING] Creating an Engine instance with the Logger and Recovery middleware already attached.

[GIN-debug] [WARNING] Running in "debug" mode. Switch to "release" mode in production.
- using env:   export GIN_MODE=release
- using code:  gin.SetMode(gin.ReleaseMode)

[GIN-debug] GET    /students          --> main.main.func1 (3 handlers)
[GIN-debug] [WARNING] You trusted all proxies, this is NOT safe. We recommend you to set a value.
Please check https://pkg.go.dev/github.com/gin-gonic/gin#readme-don-t-trust-all-proxies for details.
[GIN-debug] Listening and serving HTTP on :8080
[GIN] 2025/06/18 - 11:03:03 |+[97;42m 200 +[0m|      0s |      ::1 |+[97;44m GET      +[0m "/students"
[GIN] 2025/06/18 - 11:03:03 |+[90;43m 404 +[0m|      0s |      ::1 |+[97;44m GET      +[0m "/favicon.ico"
```

3)

```
package main
```

```
import (
    "fmt"
    "sync"
```

```

)

// Define the worker function
func worker(id int, tasks <-chan int, wg *sync.WaitGroup) {
    defer wg.Done() // Called once when the worker finishes

    for task := range tasks {
        result := task * task
        fmt.Printf("Worker %d processed task %d (result:
%d)\n", id, task, result)
    }
}

func main() {
    var wg sync.WaitGroup
    tasks := make(chan int, 6) // Buffered channel with 6 tasks

    // Start 3 workers
    for w := 1; w <= 3; w++ {
        wg.Add(1)
        go worker(w, tasks, &wg)
    }

    // Send 6 tasks (numbers to square)
    for i := 1; i <= 6; i++ {
        tasks <- i
    }

    close(tasks) // Close the channel so workers stop when all
tasks are read

    // Wait for all workers to finish
    wg.Wait()
    fmt.Println("All tasks completed.")
}

```

```
[Running] go run "c:\Users\Student\Desktop\PFAI ESE 2025\3.go"
Worker 2 processed task 2 (result: 4)
Worker 1 processed task 1 (result: 1)
Worker 1 processed task 5 (result: 25)
Worker 1 processed task 6 (result: 36)
Worker 2 processed task 4 (result: 16)
Worker 3 processed task 3 (result: 9)
All tasks completed.

[Done] exited with code=0 in 0.258 seconds
```

Part 2: JULIA

- Create a simple neural network in julia using [Flux.jl](#) that learns to predict output $y = 2x$ from inputs [1.0, 2.0 3.0]
- Include model definition, training, and display learned weight.

```
using Flux

# Define input and target data
X = [1.0, 2.0, 3.0] # Input data
y = [2.0, 4.0, 6.0] # Target outputs (y = 2x)

# Reshape data for Flux (each column is a sample)
X_train = reshape(X, 1, :) # 1x3 matrix
y_train = reshape(y, 1, :) # 1x3 matrix

# Define a simple linear model (single neuron, no bias for cleaner
y=2x learning)
model = Dense(1 => 1, identity; bias=false)

# Define loss function that takes model as parameter
function loss_fn(m)
    return Flux.mse(m(X_train), y_train)
end

# Define optimizer with higher learning rate
opt = Flux.setup(Adam(0.1), model)

# Training parameters
epochs = 1000
```

```

# Training loop
println("Training the neural network...")
println("Initial weight: ", model.weight[1])

for epoch in 1:epochs
    # Compute gradients - pass the loss function and model
    grads = Flux.gradient(loss_fn, model)

    # Update parameters
    Flux.update!(opt, model, grads[1])

    # Print progress every 100 epochs
    if epoch % 100 == 0
        current_loss = loss_fn(model)
        println("Epoch $epoch: Loss = $(round(current_loss,
digits=6)), Weight = $(round(model.weight[1], digits=4))")
    end
end

# Display final results
println("\n" * "="^50)
println("Training completed!")
println("Final learned weight: ", round(model.weight[1], digits=4))
println("Target weight: 2.0")
println("Final loss: ", round(loss_fn(model), digits=6))

# Test the model
println("\nTesting the model:")
for i in 1:length(X)
    input_val = X[i]
    predicted = model([input_val])[1]
    actual = y[i]
    println("Input: $input_val, Predicted: $(round(predicted,
digits=4)), Actual: $actual")
end

# Verify with new data point
test_input = 5.0
test_prediction = model([test_input])[1]

```



```
expected_output = 2.0 * test_input
println("\nNew test:")
println("Input: $test_input, Predicted: $(round(test_prediction,
digits=4)), Expected: $expected_output")
```

```
➔ Training the neural network...
Initial weight: -1.4703591
Epoch 100: Loss = 0.000644, Weight = 1.9883
Epoch 200: Loss = 0.0, Weight = 2.0
Epoch 300: Loss = 0.0, Weight = 2.0
Epoch 400: Loss = 0.0, Weight = 2.0
Epoch 500: Loss = 0.0, Weight = 2.0
Epoch 600: Loss = 0.0, Weight = 2.0
Epoch 700: Loss = 0.0, Weight = 2.0
Epoch 800: Loss = 0.0, Weight = 2.0
Epoch 900: Loss = 0.0, Weight = 2.0
Epoch 1000: Loss = 0.0, Weight = 2.0
```

```
=====
Training completed!
Final learned weight: 2.0
Target weight: 2.0
Final loss: 0.0
```

```
Testing the model:
Input: 1.0, Predicted: 2.0, Actual: 2.0
Input: 2.0, Predicted: 4.0, Actual: 4.0
Input: 3.0, Predicted: 6.0, Actual: 6.0
```

```
New test:
Input: 5.0, Predicted: 10.0, Expected: 10.0
```